

Programación Distribuida y Tiempo Real

Sincronización – Relojes Lógicos

Relojes Lógicos

- Se deja de lado la asociación reloj-fecha y hora
 - No se corrige *drift*, ni *offset*, ni *skew*
- No se conserva la noción de proporcionalidad de tiempo



Relojes Lógicos

- Se deja de lado la asociación reloj-fecha y hora
 - No se corrige *drift*, ni *offset*, ni *skew*
- No se conserva la noción de proporcionalidad de tiempo
- Se focaliza la relación de orden de la sucesión de eventos
- En la base de todo el desarrollo: relación “antes de”
- CS 417: Distributed Systems, Paul Krzyzanowski
<https://www.cs.rutgers.edu/~pxk/417/index.html>

Distributed Systems

Logical Clocks

Paul Krzyzanowski
pxk@cs.rutgers.edu

Except as otherwise noted, the content of this presentation is licensed under the Creative Commons Attribution 2.5 License.

Relojes Lógicos

Logical clocks

Assign sequence numbers to messages

- All cooperating processes can agree on order of events
- vs. **physical clocks**: time of day

Assume no central time source

- Each system maintains its own local clock
- No total ordering of events
 - No concept of *happened-when*

Relojes Lógicos

Happened-before

Lamport's "happened-before" notation

$a \rightarrow b$ event a happened before event b

e.g.: a : message being sent, b : message receipt

Transitive:

if $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$

Relojes Lógicos

Assign "clock" value to each event

- if $a \rightarrow b$ then $\text{clock}(a) < \text{clock}(b)$
- since time cannot run backwards

- En principio, un contador de eventos locales
 - Ante la ocurrencia de un evento
 - Incrementar el valor del contador
 - “Tiempo” de ocurrencia del evento: valor del contador
 - En cada computadora es sencillo
 - Contador + “estampilla” para los eventos
 - Seguro se respeta la relación “antes de”
-

Relojes Lógicos

Assign "clock" value to each event

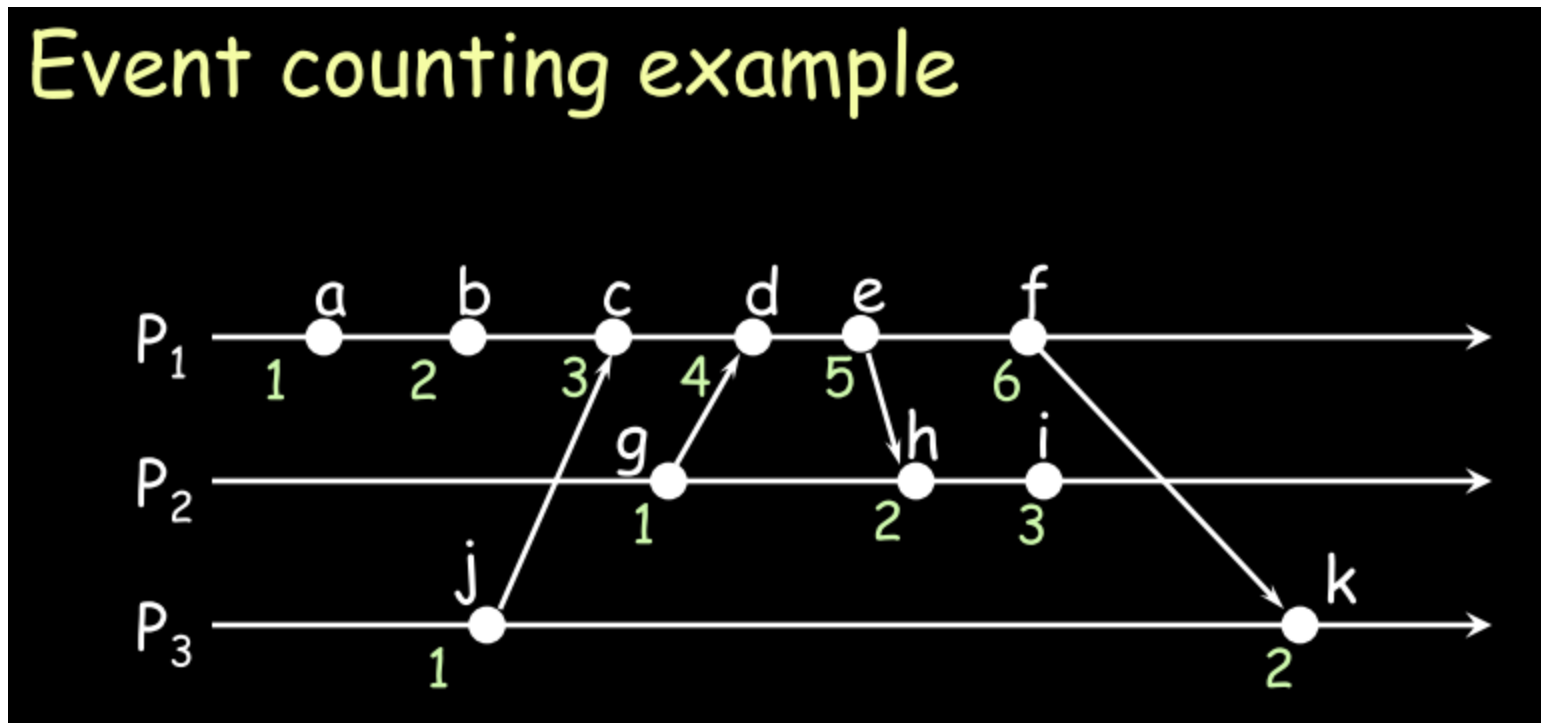
- if $a \rightarrow b$ then $\text{clock}(a) < \text{clock}(b)$
- since time cannot run backwards

- En principio, un contador de eventos locales
- Ante la ocurrencia de un evento
 - Incrementar el valor del contador
 - “Tiempo” de ocurrencia del evento: valor del contador
- En cada computadora es sencillo
 - Contador + “estampilla” para los eventos
 - Seguro se respeta la relación “antes de”

Relojes Lógicos

- Si solo contamos...

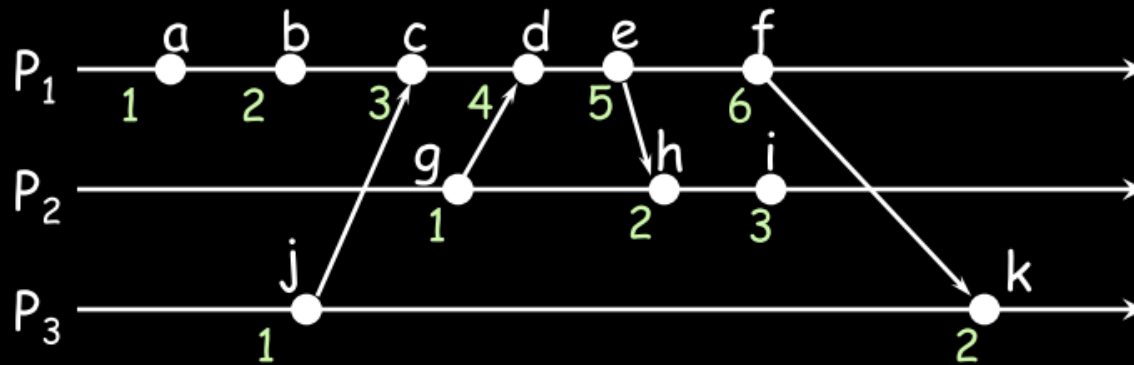
Event counting example



Relojes Lógicos

- Si solo contamos...

Event counting example



Bad ordering:

$e \rightarrow h$

$f \rightarrow k$

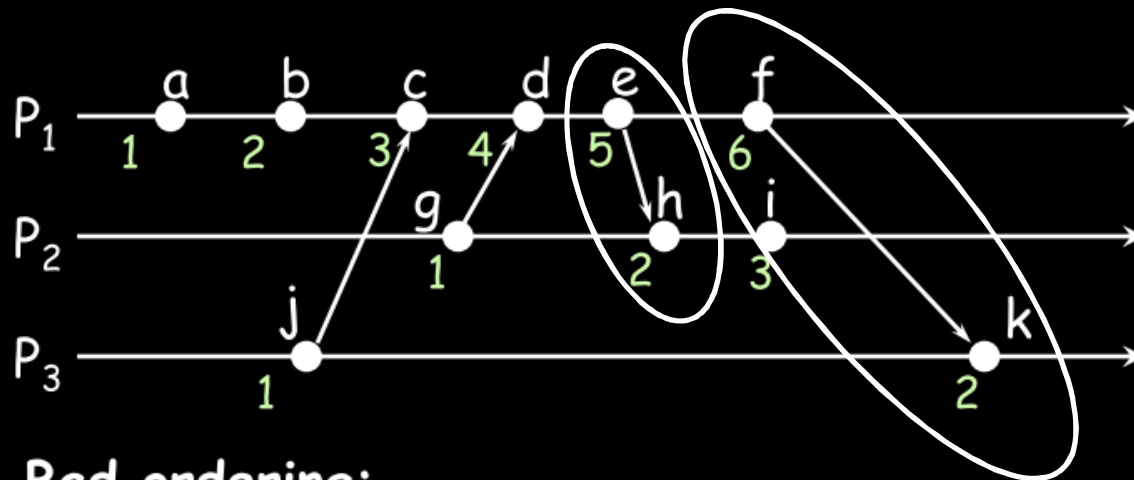
No se mantiene

- if $a \rightarrow b$ then $\text{clock}(a) < \text{clock}(b)$

Relojes Lógicos

- Si solo contamos...

Event counting example



Bad ordering:

$e \rightarrow h$

$f \rightarrow k$

No se mantiene

- if $a \rightarrow b$ then $\text{clock}(a) < \text{clock}(b)$

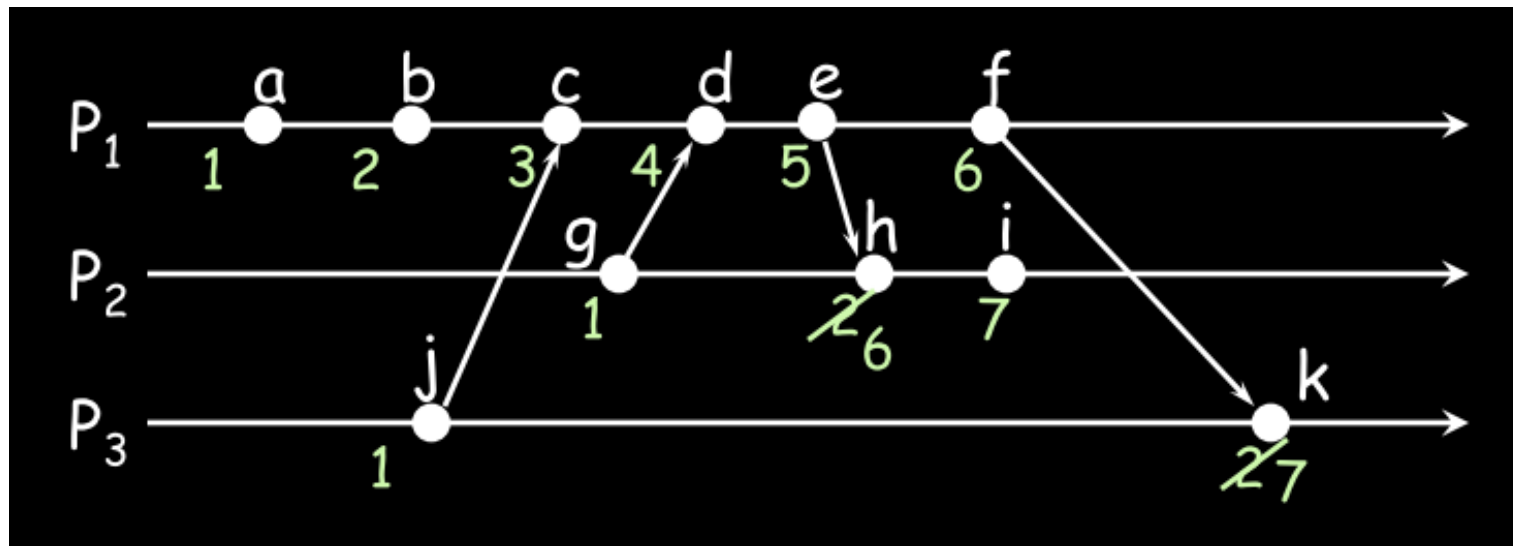
Relojes Lógicos

Lamport's algorithm

- Each message carries a timestamp of the sender's clock
- When a message arrives:
 - if receiver's clock < message timestamp
set system clock to (message timestamp + 1)
 - else do nothing
- Clock must be advanced between any two events in the same process

Relojes Lógicos

Lamport's algorithm



- Localmente se avanza
- Mensaje recibido: verificar contador de envío
- Monotónicamente creciente en todos los procesos

Relojes Lógicos

Lamport's algorithm

If a and b occur on different processes that do not exchange messages, then neither $a \rightarrow b$ nor $b \rightarrow a$ are true

- These events are **concurrent**

Algorithm allows us to maintain time ordering among related events

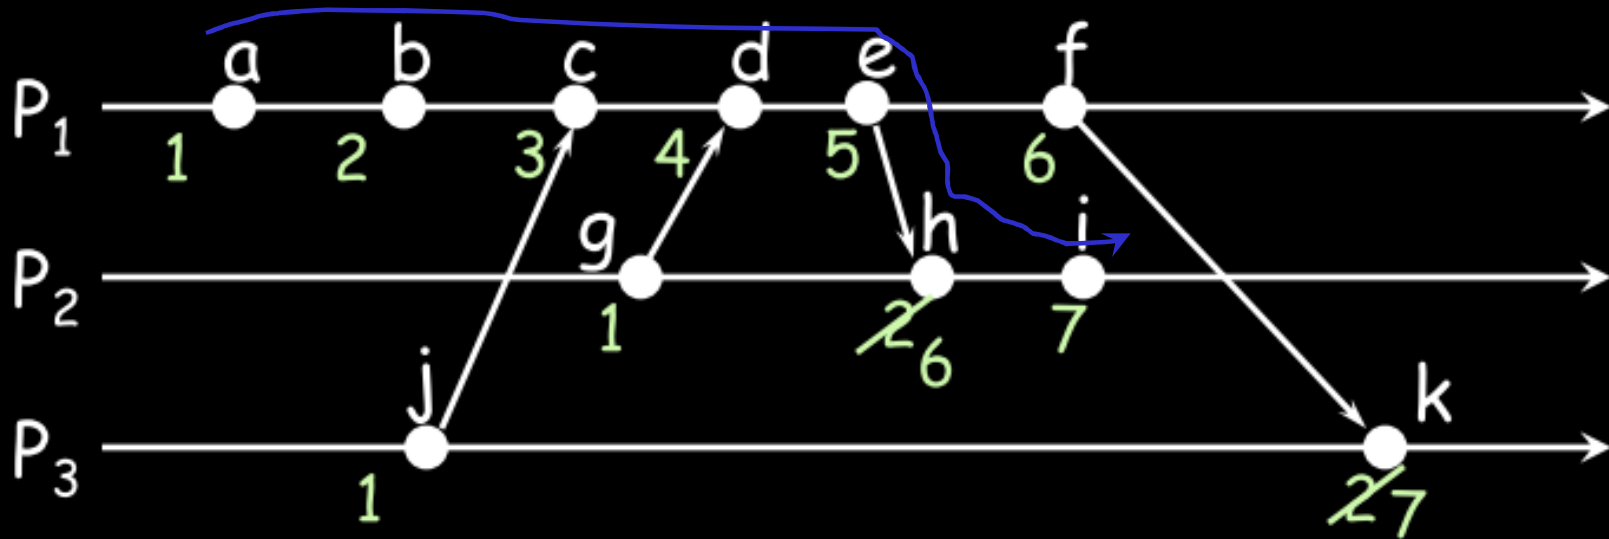
- *Partial ordering*

- Each event has a **Lamport timestamp** attached to it
- For any two events, where $a \rightarrow b$:
 $L(a) < L(b)$

Relojes Lógicos

Lamport's algorithm

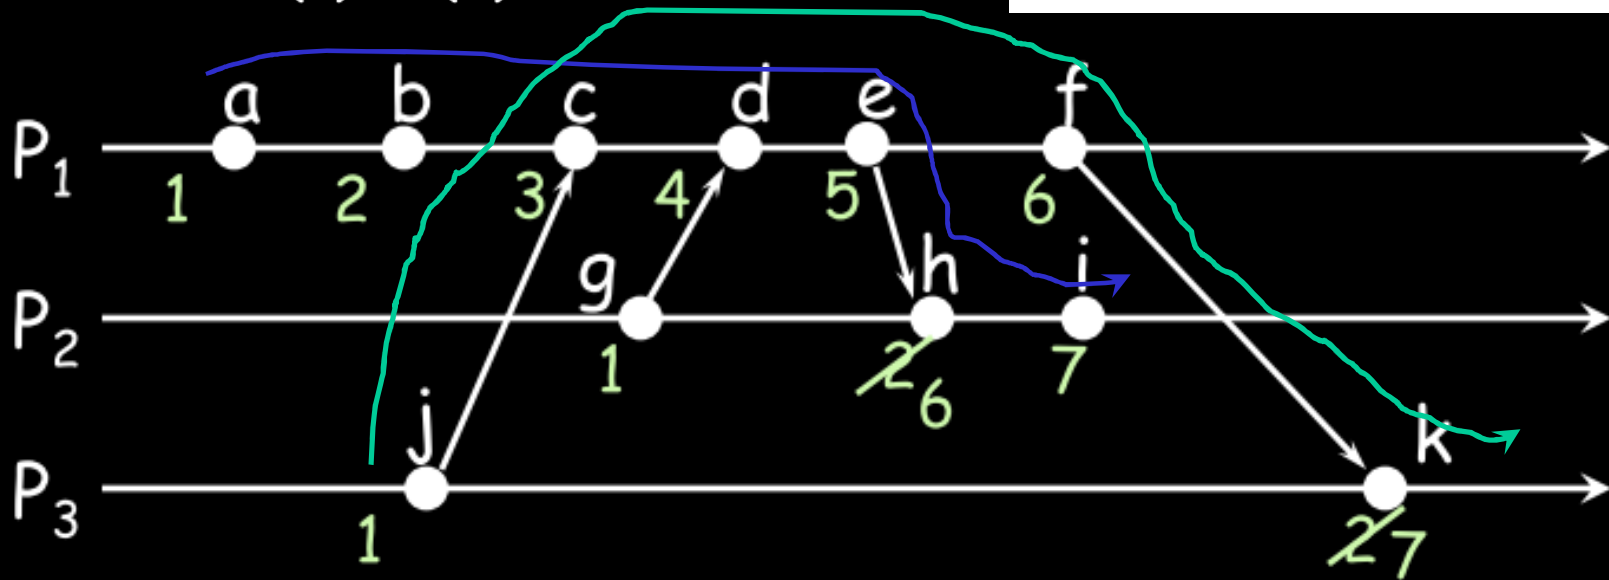
- For any two events, where $a \rightarrow b$:
 $L(a) < L(b)$



Relojes Lógicos

Lamport's algorithm

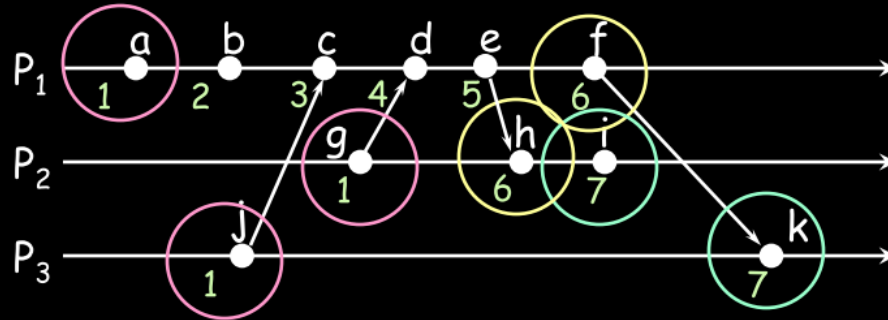
- For any two events, where $a \rightarrow b$:
 $L(a) < L(b)$



Relojes Lógicos

Lamport's algorithm

Problem: Identical timestamps



$a \rightarrow b, b \rightarrow c, \dots$: local events sequenced

$i \rightarrow c, f \rightarrow d, d \rightarrow g, \dots$: Lamport imposes a
send \rightarrow *receive* relationship

Concurrent events (e.g., a & i) may have
the same timestamp ... or not

Relojes Lógicos

Lamport's algorithm

Unique timestamps (total ordering)

We can force each timestamp to be unique

- Define global logical timestamp (T_i, i)
 - T_i represents local Lamport timestamp
 - i represents process number (globally unique)
 - E.g. (host address, process ID)
- Compare timestamps:
 - $(T_i, i) < (T_j, j)$
 - if and only if
 - $T_i < T_j$ or
 - $T_i = T_j$ and $i < j$

Does not relate to event ordering

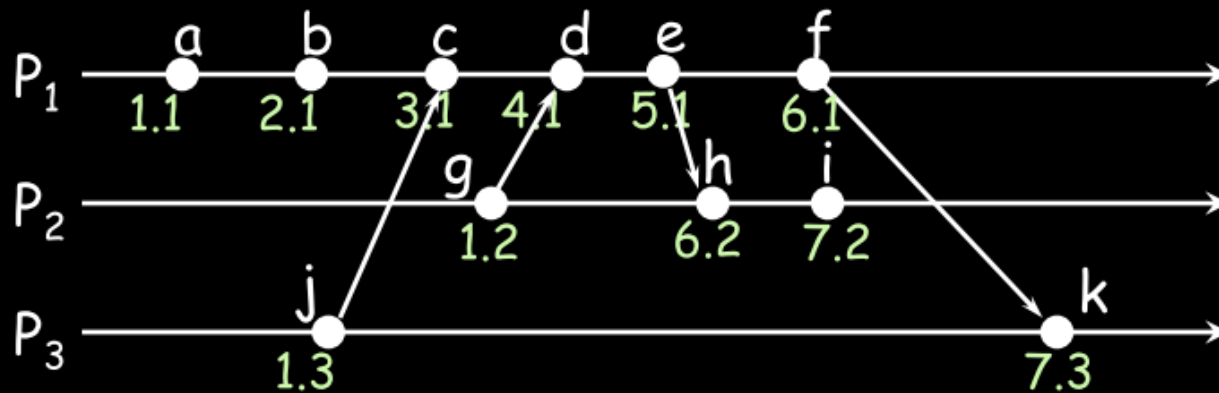
Relojes Lógicos

Lamport's algorithm

Unique timestamps (total ordering)

Does not relate to event ordering

Unique (totally ordered) timestamps



Relojes Lógicos

Lamport's algorithm

- For any two events, where $a \rightarrow b$:
 $L(a) < L(b)$

Problem: Detecting causal relations

If $L(e) < L(e')$

- Cannot conclude that $e \rightarrow e'$

Looking at Lamport timestamps

- Cannot conclude which events are causally related
-

Relojes Lógicos

Vector clocks

Rules:

1. Vector initialized to 0 at each process
 $V_i[j] = 0$ for $i, j = 1, \dots, N$
2. Process increments its element of the vector in local vector before timestamping event:
 $V_i[i] = V_i[i] + 1$
3. Message is sent from process P_i with V_i attached to it
4. When P_j receives message, compares vectors element by element and sets local vector to higher of two values
 $V_j[i] = \max(V_i[i], V_j[i])$ for $i = 1, \dots, N$

Relojes Lógicos

Comparing vector timestamps

Define

$V = V'$ iff $V[i] = V'[i]$ for $i = 1 \dots N$

$V \leq V'$ iff $V[i] \leq V'[i]$ for $i = 1 \dots N$

For any two events e, e'

if $e \rightarrow e'$ then $V(e) < V(e')$

• Just like Lamport's algorithm

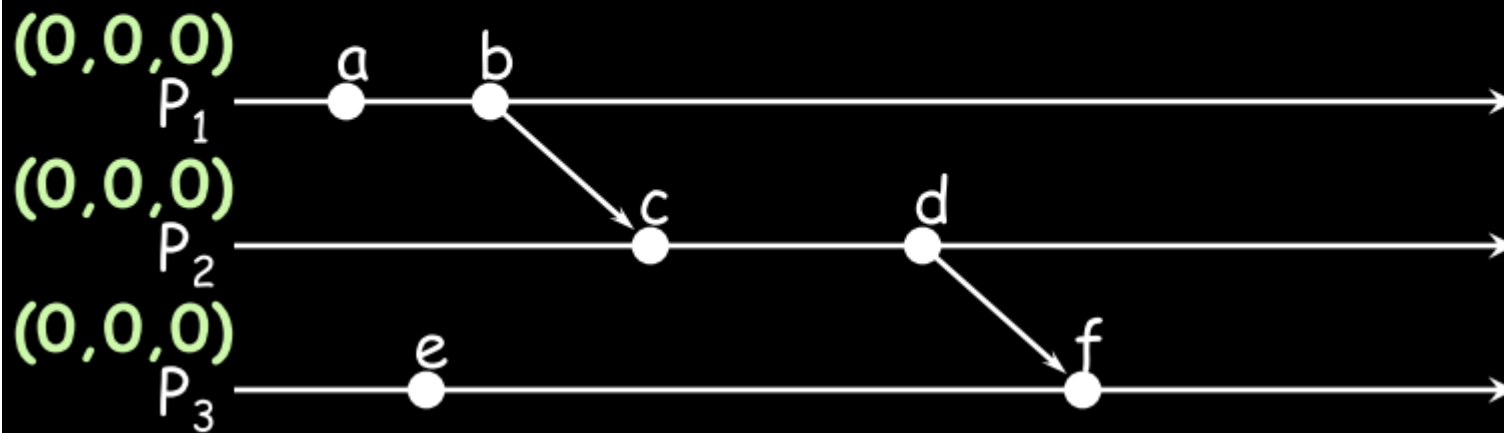
if $V(e) < V(e')$ then $e \rightarrow e'$

Two events are **concurrent** if neither

$V(e) \leq V(e')$ nor $V(e') \leq V(e)$

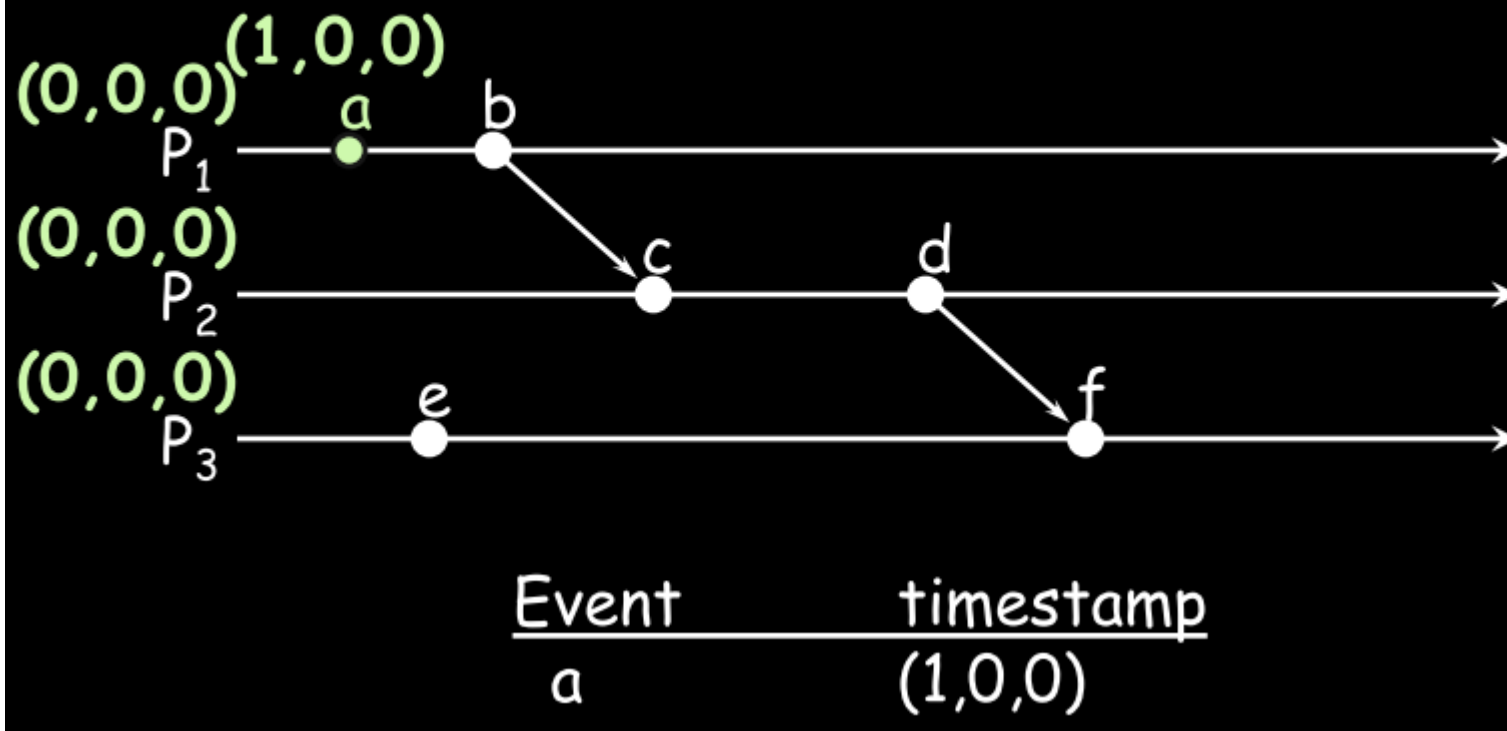
Relojes Lógicos

Vector timestamps



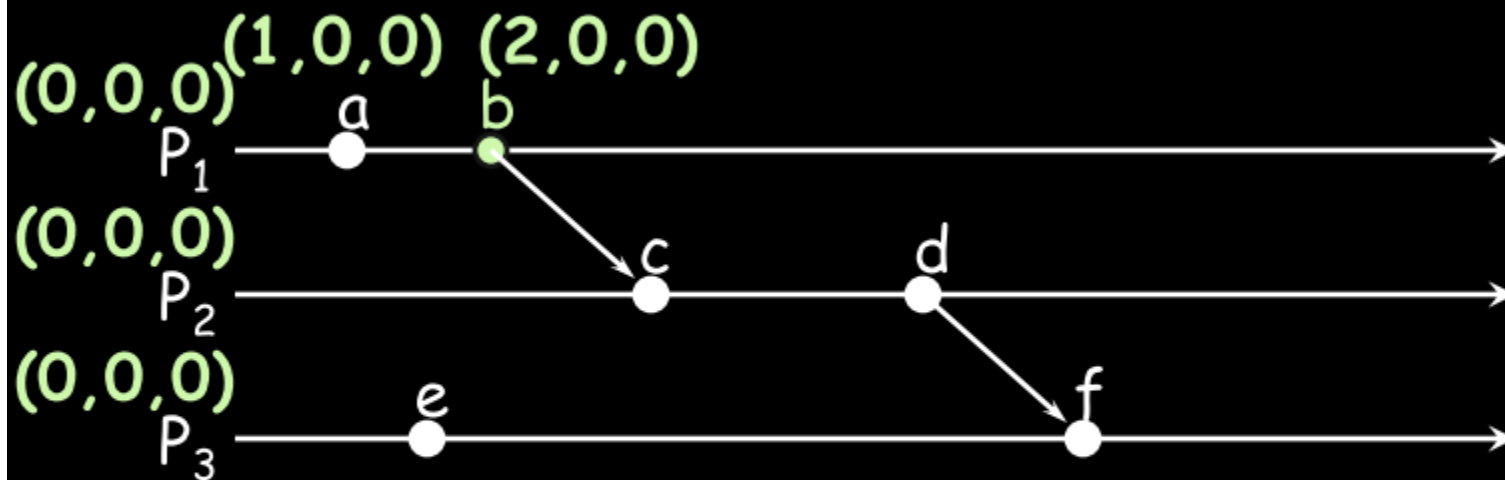
Relojes Lógicos

Vector timestamps



Relojes Lógicos

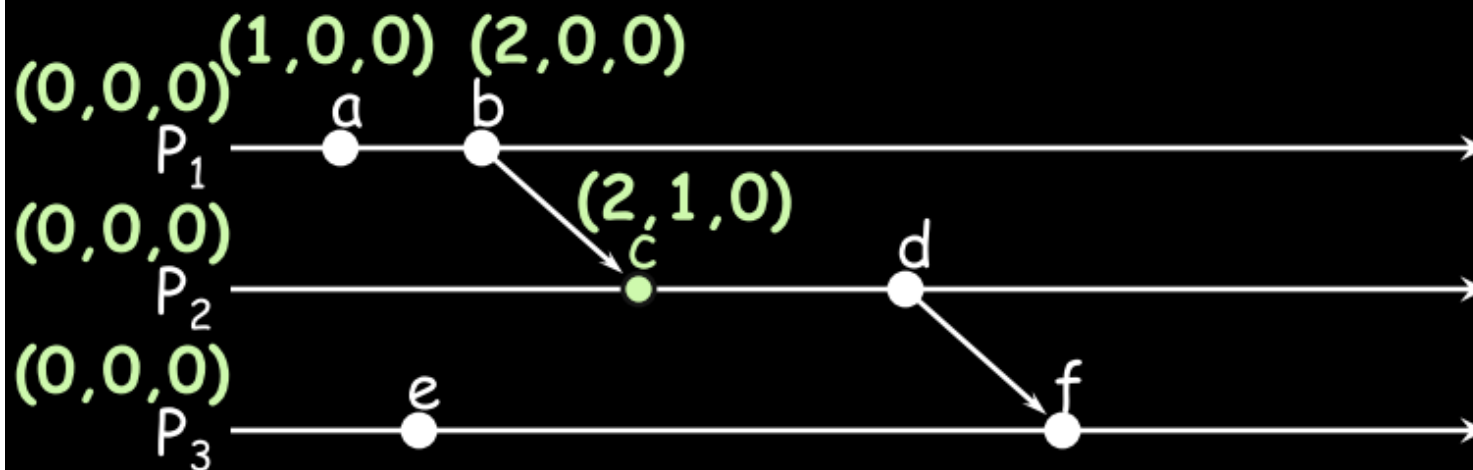
Vector timestamps



<u>Event</u>	<u>timestamp</u>
a	$(1,0,0)$
b	$(2,0,0)$

Relojes Lógicos

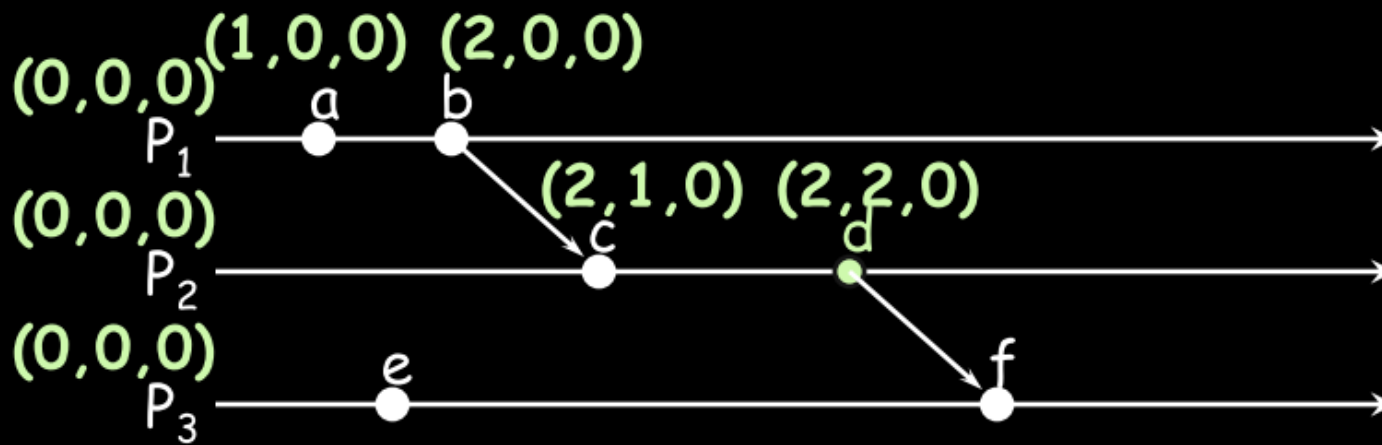
Vector timestamps



<u>Event</u>	<u>timestamp</u>
a	$(1,0,0)$
b	$(2,0,0)$
c	$(2,1,0)$

Relojes Lógicos

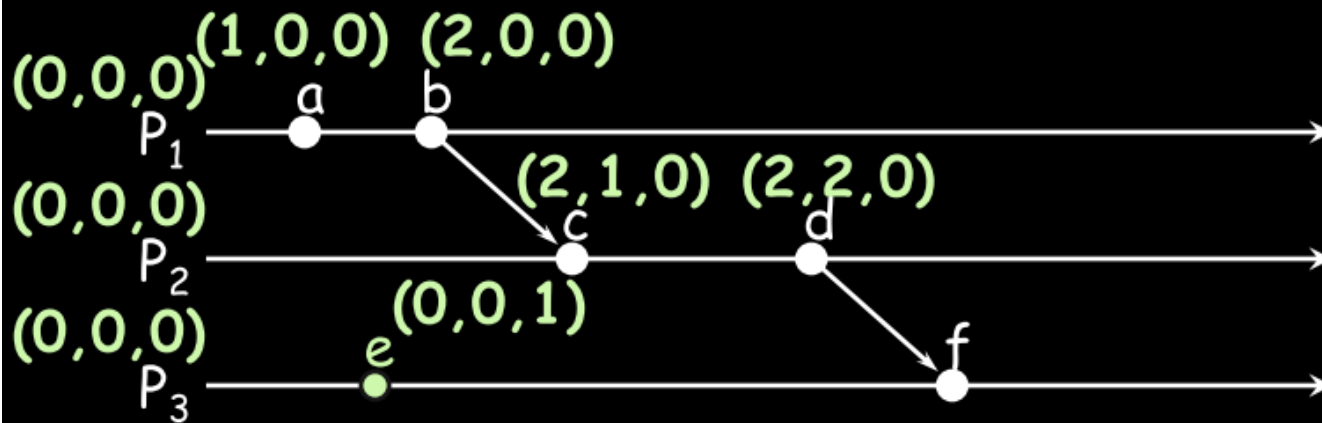
Vector timestamps



Event	timestamp
a	(1,0,0)
b	(2,0,0)
c	(2,1,0)
d	(2,2,0)

Relojes Lógicos

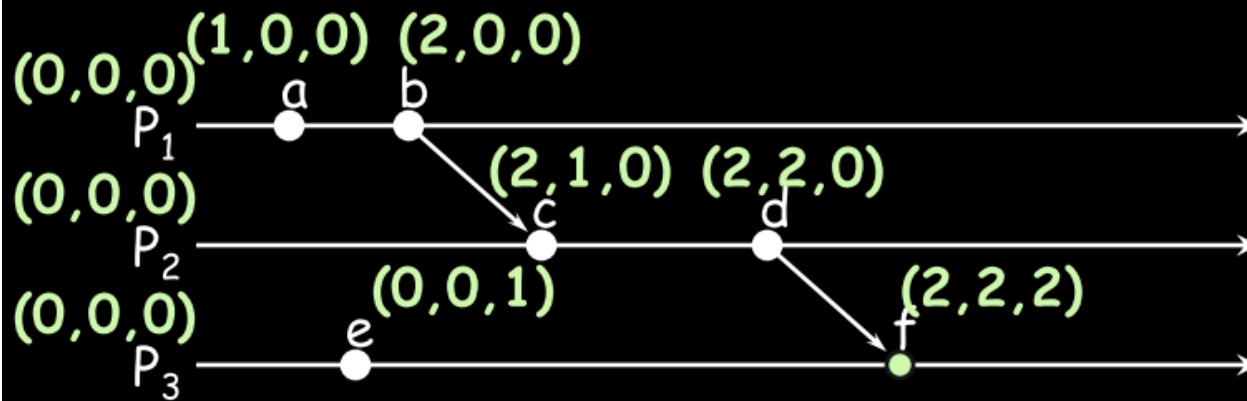
Vector timestamps



<u>Event</u>	<u>timestamp</u>
a	(1,0,0)
b	(2,0,0)
c	(2,1,0)
d	(2,2,0)
e	(0,0,1)

Relojes Lógicos

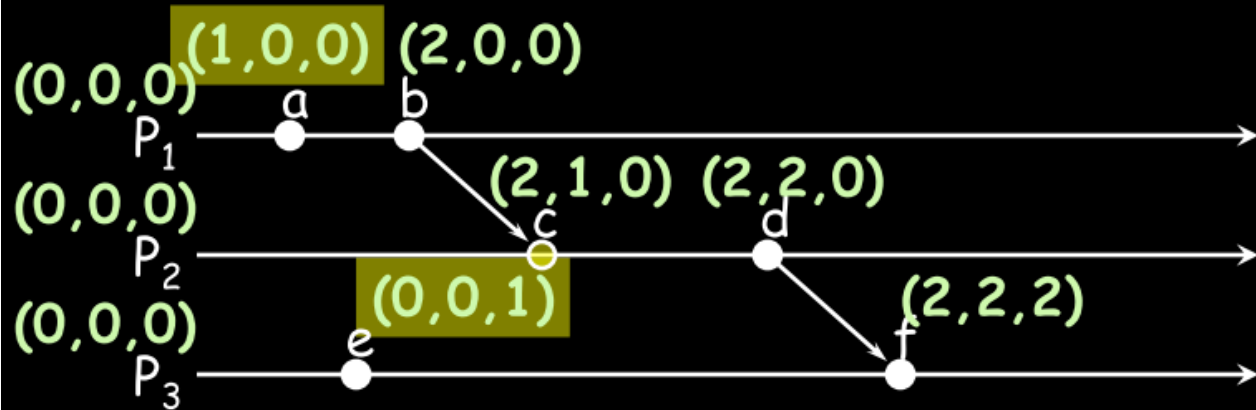
Vector timestamps



Event	timestamp
a	(1,0,0)
b	(2,0,0)
c	(2,1,0)
d	(2,2,0)
e	(0,0,1)
f	(2,2,2)

Relojes Lógicos

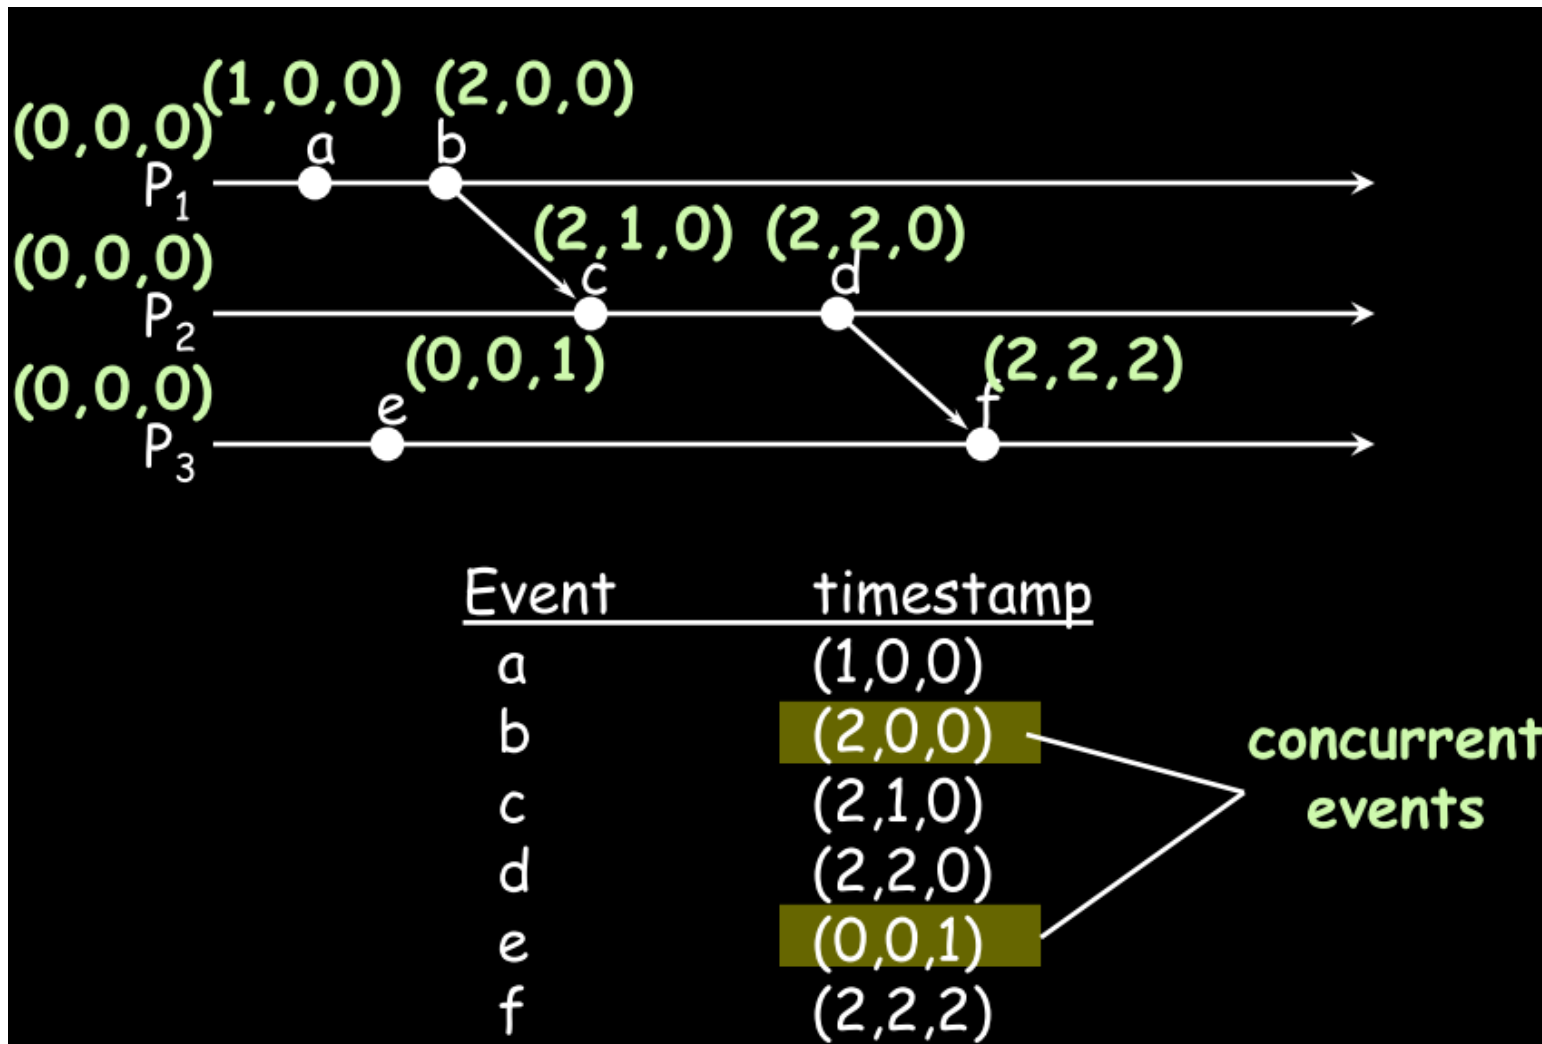
Vector timestamps



Event	timestamp
a	(1,0,0)
b	(2,0,0)
c	(2,1,0)
d	(2,2,0)
e	(0,0,1)
f	(2,2,2)

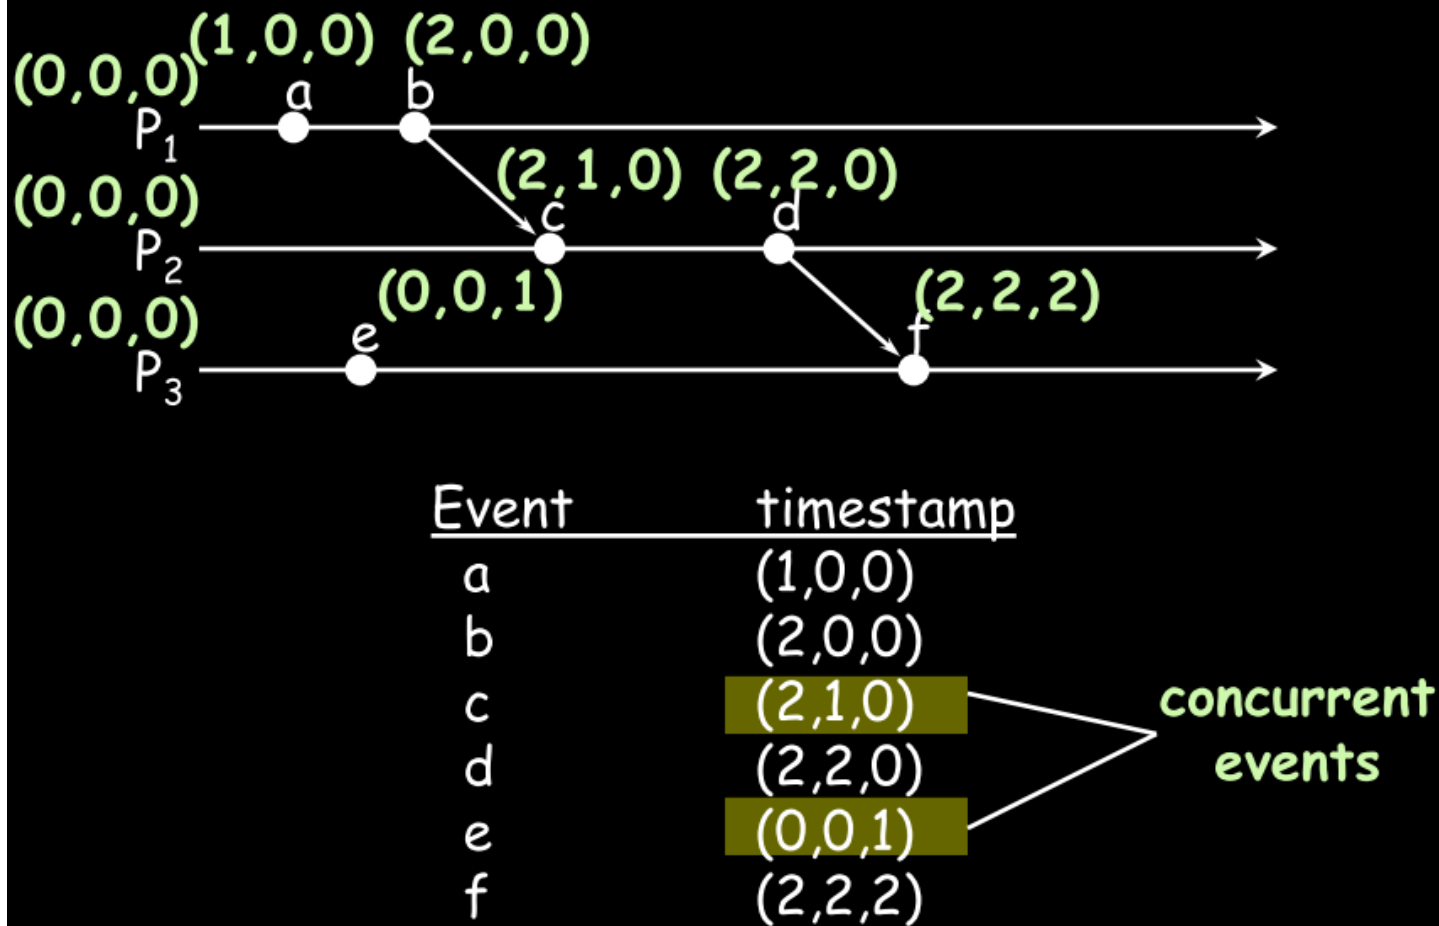
concurrent events

Relojes Lógicos



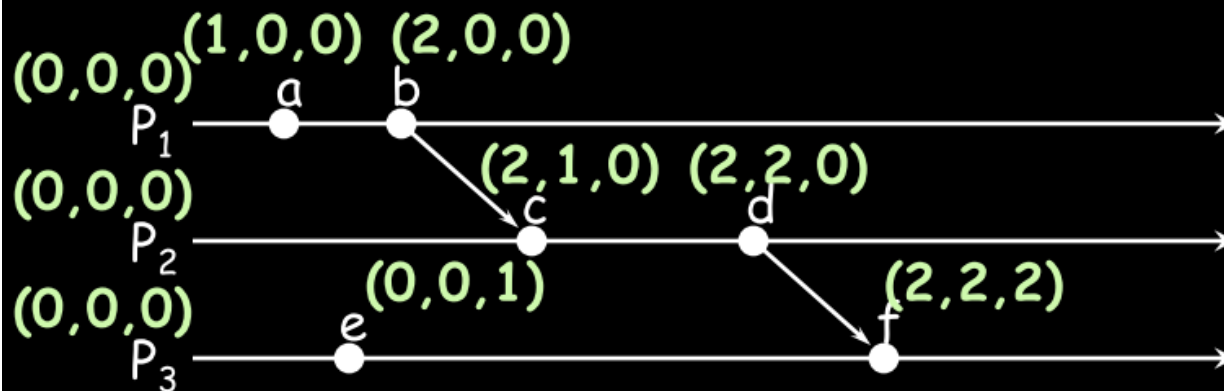
Relojes Lógicos

Vector timestamps



Relojes Lógicos

Vector timestamps



Event	timestamp
a	(1,0,0)
b	(2,0,0)
c	(2,1,0)
d	(2,2,0)
e	(0,0,1)
f	(2,2,2)

concurrent events

Relojes Lógicos

Summary: Logical Clocks & Partial Ordering

- Causality
 - If $a \rightarrow b$ then event a can affect event b
- Concurrency
 - If neither $a \rightarrow b$ nor $b \rightarrow a$ then one event cannot affect the other
- Partial Ordering
 - Causal events are sequenced
- Total Ordering
 - All events are sequenced

Dudas/Consultas

- Plataforma Ideas

